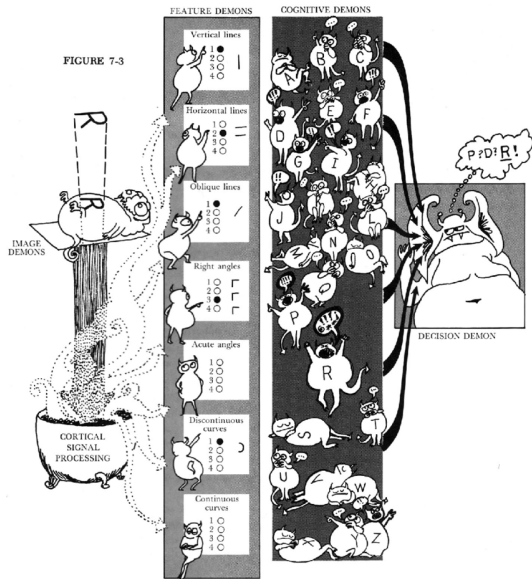


Backpropagation

Multi-layer perceptrons (MLPs)

How to solve the XOR problem ...

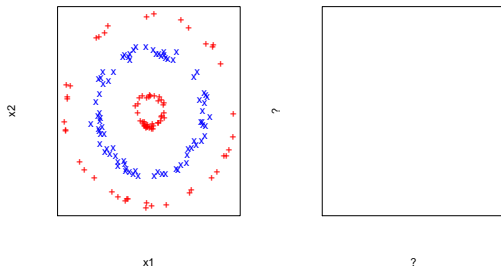
The importance of features / 1



(Lindsay and Norman's view of Selfridge's Pandemonium model, 1959).

The importance of features / 2

- Find the right features to make the task solvable:



- Engineering features by hand is hard.
- Neural networks learn features that they find important.

How many layers of features do you need?

One hidden layer is all you need **in theory** to make a “universal approximator” (Cybenko 1989; Hornik 1991).

With linear transfer functions how many layers do we need?

Neural networks and linear algebra


- Activation of a layer of neurons stored in a **vector**.
- Synapses from one layer to another stored in a **weight matrix**: \mathbf{W}_{ji} is strength of connection from unit i in one layer to unit j in the next layer.
- “The single key fact about vectors and matrices is that each vector represents a point located in space, and a matrix moves that point to a different location. Everything else is just details.” (Stone 2019, Appx. C).

Learning in multi-layer perceptrons

How to solve the **credit assignment problem**? i.e. what is the “delta” for hidden units, that have no desired output?

$$\delta_j = g'(h_j) \sum_i w_{ij} \delta_i$$

h_j is total input to unit j ; g' is 1st derivative of activation function.
See backprop.pdf and the accompanying video.



No guarantee (unlike PCT) that this will converge due to local minima.

Application: solving XOR...

DEMO in live class.

How to assess for over fitting vs under fitting

- **Underfitting:** can network solve problem?
- **Overfitting:** network too focused on learning (perhaps by rote) the training examples.
- **Generalisation:** how well does network perform on inputs not seen during learning?
- Where is the “Goldilocks spot”?
- Run on **validation set** during learning. Plot error as a function of training time (epochs). Assess after on test set.



- Other approaches (k-fold validation) also feasible. Be careful about time-dependencies!

Some terms

1. **Online learning:** learn after every input. (each **iteration**).
2. **Batch learning:** wait until all training samples have been presented (each **epoch**).
3. **Mini batch:** break data into several groups (make sure each is balanced).

Applications: family trees (Hinton 1986; Paccanaro and Hinton, 2000).

- How to predict family trees? (person X) (relationship) (who?)
- e.g. (Charlotte) (has aunt) (who) \Rightarrow Jennifer or Margaret
- 2 family trees of 12 people = 24 people.
- 12 possible relationships:
husband, wife, son, daughter, father, mother
brother, sister, nephew, neice, uncle, aunt
- 104 relationships; 100 used in training; 4 for testing.


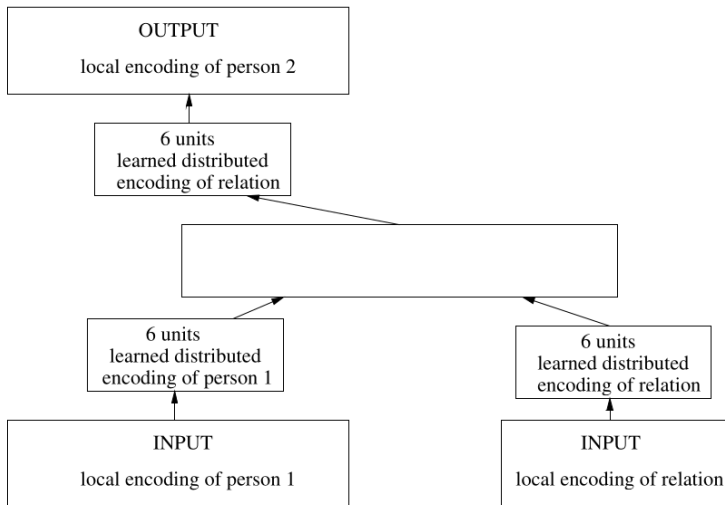


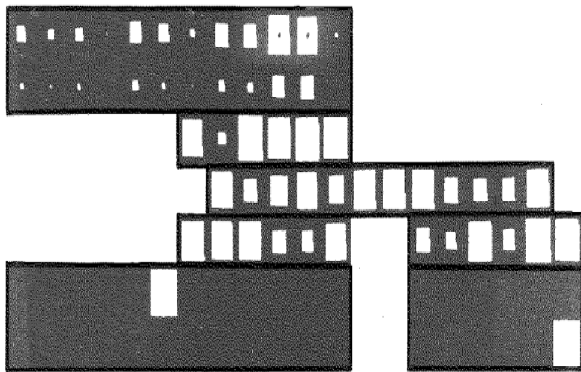
Figure 2: Two isomorphic family trees. The symbol “=” means “married to”

Family tree architecture

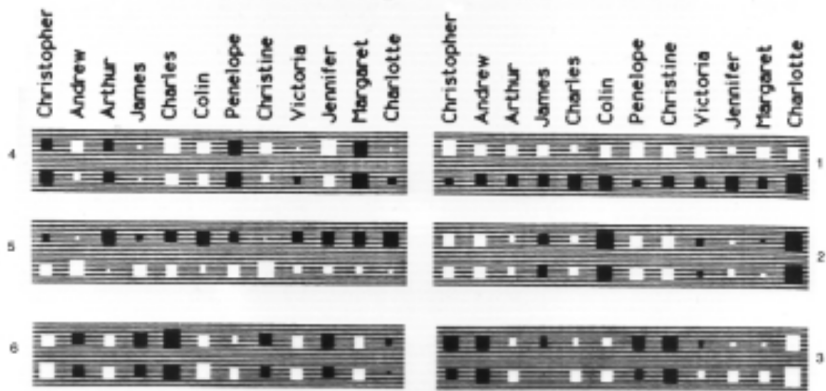


After training

When tested on: (Colin) (has Aunt). See next slide for interpretation of outputs.



What are the hidden units doing?



Generalisation

Very small data set, but normally got at least 2 (of 4) test set examples correct.

Applications: NET TALK

NetTalk: Sejnowski and Rosenberg (1987)

How do humans learn to pronounce words? Develop mapping from letters to phonemes/stress.

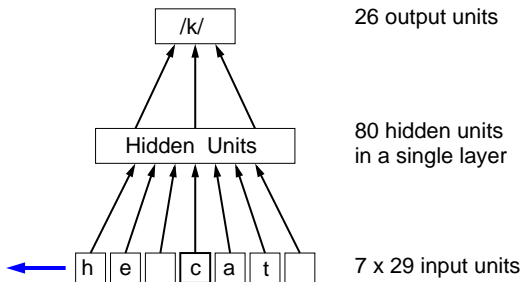
Corpus chosen from either (a) 1000 most common words from Pocket Dictionary (b) young child's informal speech.

Context of letter important: mapping from letter to pronunciation dependent on context:

phoneme	word	articulation features
/A/	bite	medium, tensed, front2+central1
/I/	bit	high, front1

Too many words (even with approx 18,000 weights) to just produce lookup table – must extract principles.

NetTalk: architecture



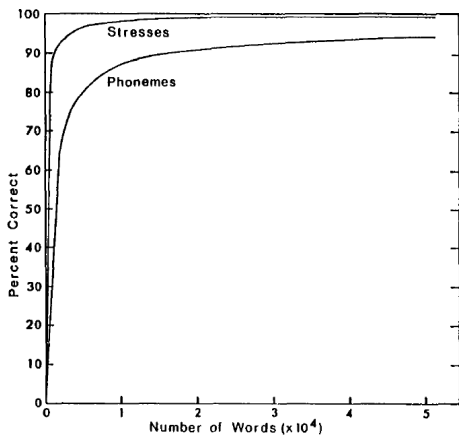


Input: each letter encoded using 1 of 29 units (26 + 3 for punctuation). Central letter is processed, using surrounding 3 letters each side to provide context.

Output: distributed representation across 21 features including vowel height, position in mouth; 5 features for stress.

NetTalk: performance

- Stage-like progression of behaviour:
 1. distinction between vowels and consonants: but same vowel for all vowels, and same consonant for all consonants. (babbling)
 2. recognition of word boundaries.
 3. pseudo words
 4. good performance (90%)
- Power-law like learning (similar to humans)
- Robust to weight damage; rapid recovery.
- Generalisation: $\approx 80\%$



NetTalk: demo

Demonstration of NetTalk:

<http://cnl.salk.edu/Media/nettalk.mp3>

Part 1 : learning from zero weights (0:37).

Part 2 : learning after 10,000 iterations (3:18).

Part 3 : performance on unseen text during learning (5:02).

NB: Similar architecture of moving input window used to predict protein structure (α helix, β sheet, other) of central part of 13 amino acids (Qian & Sejnowski, 1988).

Summary

1. We have now seen the work of two of the leading pioneers in the field: Terry Sejnowski and Geoff Hinton.
2. Also known as “the Lennon and McCartney of neural networks” (Jim Stone).