

Boltzmann machine

Move from a deterministic to stochastic regime for asynchronous update:

$$\text{total input: } a_i = \sum_{j=1}^N w_{ij} v_j$$

$$\text{update rule: } P(v_i = 1) = f(a_i)$$

$$\text{where } f(a_i) = \frac{1}{1 + \exp(-a_i)}$$

System converges to an equilibrium state for the states \mathbf{v} given by:

$$\text{energy function: } E(\mathbf{v}) = -\frac{1}{2} \mathbf{v}^T W \mathbf{v}$$

$$\text{Boltzmann distribution: } P(\mathbf{v}) = \frac{\exp(-E(\mathbf{v}))}{\sum_{\mathbf{v}} \exp(-E(\mathbf{v}))}$$

Can also introduce “hidden units” to detect higher order correlations (not just pairwise).

1/11

2/11

Restricted Boltzmann Machine (RBM)

Two layer network, with input layer connected to/from hidden layer; no within-layer connections.

In a **wake** phase, input units are clamped on, and drive hidden layer. In a **sleep** phase, hidden layer units can drive inputs.

Trained using a procedure called Contrastive Divergence (Stone, Chapter 7).

Much more efficient than simulated annealing for Boltzmann machines.

Stacked RBMs

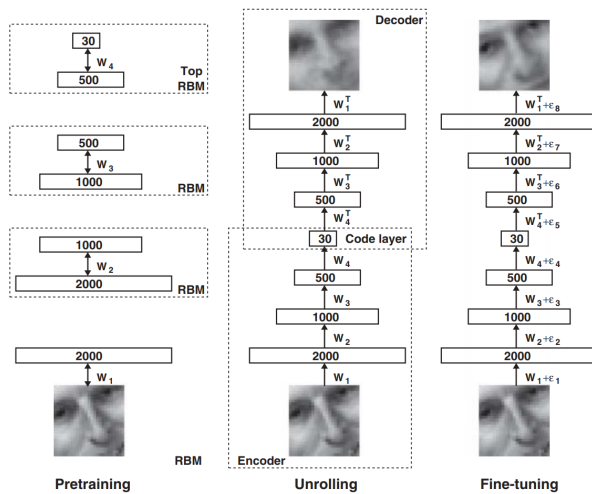
Can train a stack of RBMs one-by-one, such that a hidden layer, once trained is used as input layer to next RBM.

3/11

4/11


Autoencoders (Hinton and Salakhutdinov 2006)

After training stacked RBM, we have an encoding network, which can be “flipped” to make a decoder with same weights. Can then refine whole net with backprop.



5/11

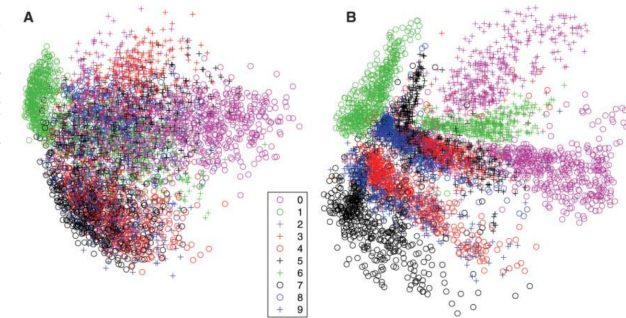
Variational autoencoders



7/11

MNIST visualisation (Figure 3 of Hinton and Salakhutdinov 2006)

Fig. 3. (A) The two-dimensional codes for 500 digits of each class produced by taking the first two principal components of all 60,000 training images. (B) The two-dimensional codes found by a 784-1000-500-250-2 autoencoder. For an alternative visualization, see (8).



Sup layer above autoencoder classified MNIST with 1.6% error. (Stone, p96). Netflix 1 million USD prize won by team using SVD + RBMs; not used as films moved to online delivery.
<https://www.techdirt.com/articles/20120409/03412518422/why-netflix-never-implemented-algorithm-that-won-netflix-1-million-prize>

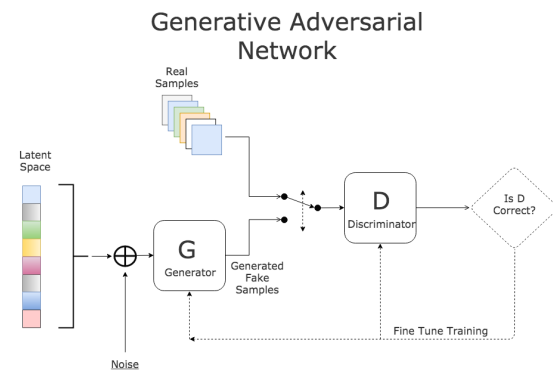
6/11

Sampling the latent space (Goodfellow, Figure 20.6)



8/11

Generative adversarial networks (Goodfellow et al 2014)



Source: <https://www.kdnuggets.com/2017/01/generative-adversarial-networks-hot-topic-machine-learning.html>

For further information:

<http://bamos.github.io/2016/08/09/deep-completion/>

1. **Discriminator** spots real vs fake training samples. Adjust weights to increase discrimination.
2. **Generator** adjusts weights to generate images that are more likely to be classified as training images.

Radford et al. (2015), figure 4


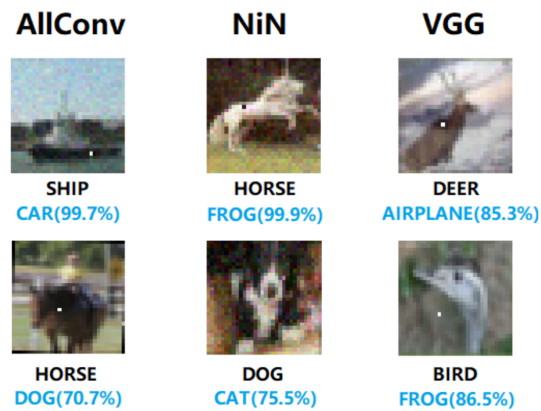


Figure 4: Top rows: Interpolation between a series of 9 random points in Z show that the space learned has smooth transitions, with every image in the space plausibly looking like a bedroom. In the 6th row, you see a room without a window slowly transforming into a room with a giant window. In the 10th row, you see what appears to be a TV slowly being transformed into a window.

9 / 11

10 / 11

Fooling Deep Networks with adversarial samples



Su et al (2017)

See also <https://arxiv.org/pdf/1707.08945.pdf> for robust attacks on stop signs.